

Diffusion models

CS 4804: Introduction to AI

Fall 2025

<https://tuvllms.github.io/ai-fall-2025/>

Tu Vu



based on Issac Ke's lecture

Logistics

- HW 2 **due today**
- Final presentations: **12/4 & 12/9**
 - Sign-up form available on Piazza

Grok-4.1


xAI's Grok 4.1 (thinking) & Grok 4.1 rank #1 and #2 of the Text Arena

Rank ↕	Rank Spread ○ (Upper-Lower)	Model ↕	Score ↓	95% CI (±) ↕	Votes ↕	Organization ↕	License ↕
1	1 ↔ 2	 grok-4.1-thinking	1483	 Preliminary ±11	3,298	xAI	Proprietary
2	1 ↔ 4	 grok-4.1	1465	 Preliminary ±11	3,413	xAI	Proprietary
3	2 ↔ 6	 gemini-2.5-pro	1452	±4	65,503	Google	Proprietary
4	2 ↔ 8	 claude-sonnet-4-5-20250929-thinking-32k	1450	±5	16,549	Anthropic	Proprietary
5	3 ↔ 7	 claude-opus-4-1-20250805-thinking-16k	1449	±4	32,080	Anthropic	Proprietary
6	3 ↔ 11	 claude-sonnet-4-5-20250929	1445	±6	11,121	Anthropic	Proprietary
7	4 ↔ 13	 gpt-4.5-preview-2025-02-27	1442	±6	14,644	OpenAI	Proprietary
8	5 ↔ 13	 claude-opus-4-1-20250805	1440	±4	44,792	Anthropic	Proprietary
9	6 ↔ 13	 chatgpt-4o-latest-20250326	1438	±4	51,321	OpenAI	Proprietary
10	6 ↔ 14	 gpt-5-high	1437	±5	32,955	OpenAI	Proprietary
11	7 ↔ 14	 o3-2025-04-16	1434	±4	61,685	OpenAI	Proprietary
12	7 ↔ 16	 qwen3-max-preview	1434	±5	28,191	Alibaba	Proprietary

Gemini 3.0

Benchmark	Description		Gemini 3 Pro	Gemini 2.5 Pro	Claude Sonnet 4.5	GPT-5.1
Humanity's Last Exam	Academic reasoning	No tools	37.5%	21.6%	13.7%	26.5%
ARC-AGI-2	Visual reasoning puzzles	ARC Prize Verified	31.1%	4.9%	13.6%	17.6%
GPQA Diamond	Scientific knowledge	No tools	91.9%	86.4%	83.4%	88.1%
AIME 2025	Mathematics	No tools	95.0%	88.0%	87.0%	94.0%
		With code execution	100%	—	100%	—
MathArena Apex	Challenging Math Contest problems		23.4%	0.5%	1.6%	1.0%
MMMU-Pro	Multimodal understanding and reasoning		81.0%	68.0%	68.0%	80.8%
ScreenSpot-Pro	Screen understanding		72.7%	11.4%	36.2%	3.5%
CharXiv Reasoning	Information synthesis from complex charts		81.4%	69.6%	68.5%	69.5%
OmniDocBench 1.5	OCR	Overall Edit Distance, lower is better	0.115	0.145	0.145	0.147
Video-MMMU	Knowledge acquisition from videos		87.6%	83.6%	77.8%	80.4%
LiveCodeBench Pro	Competitive coding problems from Codeforces, ICPC, and IOI	Elo Rating, higher is better	2,439	1,775	1,418	2,243
Terminal-Bench 2.0	Agentic terminal coding	Terminus-2 agent	54.2%	32.6%	42.8%	47.6%
SWE-Bench Verified	Agentic coding	Single attempt	76.2%	59.6%	77.2%	76.3%
t2-bench	Agentic tool use		85.4%	54.9%	84.7%	80.2%
Vending-Bench 2	Long-horizon agentic tasks	Net worth (mean), higher is better	\$5,478.16	\$573.64	\$3,838.74	\$1,473.43
FACTS Benchmark Suite	Held out internal grounding, parametric, MM, and search retrieval benchmarks		70.5%	63.4%	50.4%	50.8%
SimpleQA Verified	Parametric knowledge		72.1%	54.5%	29.3%	34.9%
MMMLU	Multilingual Q&A		91.8%	89.5%	89.1%	91.0%
Global PIQA	Commonsense reasoning across 100 Languages and Cultures		93.4%	91.5%	90.1%	90.9%
MRCR v2 (8-needle)	Long context performance	128k (average)	77.0%	58.0%	47.1%	61.6%
		1M (pointwise)	26.3%	16.4%	not supported	not supported

An AI Scientist for Autonomous Discovery



Diffusion process

The laws of physics tell us that dye particles continue to move through water until their concentration becomes uniform everywhere in the beaker, which marks equilibrium.



Diffusion models power AI image/video tools



A turtle wearing sunglasses playing basketball



Image created by ChatGPT

What are diffusion models?

- A diffusion model is a generative model that first corrupts data by gradually adding noise and then learns to reverse that process so it can create new samples by iteratively denoising

Denoising Diffusion Probabilistic Models

Jonathan Ho
UC Berkeley
jonathanho@berkeley.edu

Ajay Jain
UC Berkeley
ajayj@berkeley.edu

Pieter Abbeel
UC Berkeley
pabbeel@cs.berkeley.edu


Abstract

We present high quality image synthesis results using diffusion probabilistic models, a class of latent variable models inspired by considerations from nonequilibrium thermodynamics. Our best results are obtained by training on a weighted variational bound designed according to a novel connection between diffusion probabilistic models and denoising score matching with Langevin dynamics, and our models naturally admit a progressive lossy decompression scheme that can be interpreted as a generalization of autoregressive decoding. On the unconditional CIFAR10 dataset, we obtain an Inception score of 9.46 and a state-of-the-art FID score of 3.17. On 256x256 LSUN, we obtain sample quality similar to ProgressiveGAN. Our implementation is available at <https://github.com/hojonathanho/diffusion>.

Three important concepts


- **Forward diffusion**
- **Reverse diffusion**
- **Conditional diffusion**

Forward diffusion



<https://newsletter.theaiedge.io/p/diffusion-models-stable-diffusion>

Forward diffusion (cont'd)



Grayscale images



0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

Color images




0	3	2	5	1	7	6	0	8
0	3	2	5	1	7	6	0	8
0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0


channel x height x width

Channels are usually RGB: Red, Green, and Blue

Forward diffusion (cont'd)



$(255, 0, 0) + (-2, 2, 0)$
 $\rightarrow (253, 2, 0)$
// less red + more green



When you draw a sample from a Gaussian distribution $N(\mu, \sigma^2)$, the sampled value is

$$x = \mu + \sigma\varepsilon$$

where

- x is the sampled value.
- μ is the mean that centers the distribution.
- σ is the standard deviation that determines the spread.
- ε is a standard normal variable that follows $N(0, 1)$ and provides randomness.

$$x = \mu + \sqrt{\beta} \varepsilon$$

where

- x is the sampled value.
- μ is the mean that centers the distribution.
- β is the variance.
- $\sqrt{\beta}$ is the standard deviation.
- ε is a standard normal variable that follows $N(0, 1)$ and provides randomness.

Markov chain

$$q(x_t \mid x_{t-1}) = \boxed{?} x_{t-1} + \sqrt{\beta} \epsilon$$

Objective

$$q(x_t \mid x_0) \xrightarrow[t \rightarrow \infty]{} \mathcal{N}(0, 1)$$

Noise scheduler (cont'd)

$$q(x_t \mid x_{t-1}) = \sqrt{1 - \beta} x_{t-1} + \sqrt{\beta} \epsilon$$

Noise scheduler (cont'd) Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$

$$\begin{aligned} \mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\ &= \sqrt{\alpha_t} \left(\sqrt{\alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2}^* \right) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\alpha_t - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2}^* + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1}^* \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\sqrt{\alpha_t - \alpha_t \alpha_{t-1}}^2 + \sqrt{1 - \alpha_t}^2} \boldsymbol{\epsilon}_{t-2} \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{\alpha_t - \alpha_t \alpha_{t-1} + 1 - \alpha_t} \boldsymbol{\epsilon}_{t-2} \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \boldsymbol{\epsilon}_{t-2} \\ &= \dots \\ &= \sqrt{\prod_{i=1}^t \alpha_i} \mathbf{x}_0 + \sqrt{1 - \prod_{i=1}^t \alpha_i} \boldsymbol{\epsilon}_0 \\ &= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}_0 \end{aligned}$$

Noise scheduler

Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$


$$\begin{aligned}\mathbf{x}_t &= \sqrt{\alpha_t} \mathbf{x}_{t-1} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_{t-1} && \text{;where } \boldsymbol{\epsilon}_{t-1}, \boldsymbol{\epsilon}_{t-2}, \dots \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \\ &= \sqrt{\alpha_t \alpha_{t-1}} \mathbf{x}_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\boldsymbol{\epsilon}}_{t-2} && \text{;where } \bar{\boldsymbol{\epsilon}}_{t-2} \text{ merges two Gaussians (*).} \\ &= \dots \\ &= \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}\end{aligned}$$

(*) Recall that when we merge two Gaussians with different variance, $\mathcal{N}(\mathbf{0}, \sigma_1^2 \mathbf{I})$ and $\mathcal{N}(\mathbf{0}, \sigma_2^2 \mathbf{I})$, the new distribution is $\mathcal{N}(\mathbf{0}, (\sigma_1^2 + \sigma_2^2) \mathbf{I})$. Here the merged standard deviation is $\sqrt{(1 - \alpha_t) + \alpha_t(1 - \alpha_{t-1})} = \sqrt{1 - \alpha_t \alpha_{t-1}}$.

Usually, we can afford a larger update step when the sample gets noisier, so $\beta_1 < \beta_2 < \dots < \beta_T$ and therefore $\bar{\alpha}_1 > \dots > \bar{\alpha}_T$.

<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

Reverse diffusion




<https://newsletter.theaiedge.io/p/diffusion-models-stable-diffusion>


Reverse diffusion (cont'd)

*Every block of stone has a statue
inside it and it is the task of the
sculptor to discover it.*


- Michelangelo



Reverse diffusion (cont'd)




Reverse diffusion (cont'd)




$X_{t=T}$

-




predicted
noise

=



$X_{t=T-1}$


Reverse diffusion (cont'd)




Mean squared error

$$\left(\begin{array}{c} \text{ground truth} \\ \text{noise} \end{array} - \begin{array}{c} \text{predicted} \\ \text{noise} \end{array} \right)^2$$

Reverse diffusion (cont'd)



Conditional diffusion



Text embedding of
“A turtle wearing
sunglasses playing
basketball”


The model learns the relationship between the meaning of words and how they correlate with certain denoising sequences that gradually reveal different features and shapes and edges in the picture.

Methods for incorporating text embeddings


- SAG: Self-attention guidance
 - Forces the model to pay attention to how specific portions of the prompt influenced the generation of certain regions or areas of the image
- Classifier-free diffusion guidance
 - Helps amplify the effect that certain words in the prompt on how the image is generated

OpenAI's CLIP


1. Contrastive pre-training




Stability AI's Stable Diffusion



DALL-E



Imagen



Sora

- <https://www.youtube.com/watch?v=fG3IE9dkyKY>

Thank you!